

```
# Cell 1: Setup & imports
import os
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, models, callbacks
print("TensorFlow version:", tf.__version__)

# Check GPU
device_name = tf.test.gpu_device_name()
print("GPU device:", device_name if device_name else "No GPU detected. Enable GPU in Runtime > Change runtime type")
```

TensorFlow version: 2.19.0
GPU device: /device:GPU:0

```
# Cell 2: Load CIFAR-10
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
print("Shapes:", x_train.shape, y_train.shape, x_test.shape, y_test.shape)

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
# Show a few images
plt.figure(figsize=(10,4))
for i in range(10):
    plt.subplot(2,5,i+1)
    plt.imshow(x_train[i])
    plt.title(class_names[int(y_train[i])])
    plt.axis('off')
plt.tight_layout()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 4s 0us/step
Shapes: (50000, 32, 32, 3) (50000, 1) (10000, 32, 32, 3) (10000, 1)
/tmp/ipython-input-2797491883.py:11: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will



```
# Cell 3: Preprocess
# Normalize pixel values to [0,1]
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# If using sparse_categorical_crossentropy we can keep y as integer labels
y_train_flat = y_train.flatten()
y_test_flat = y_test.flatten()
```

```
# Cell 4: Data augmentation using tf.keras preprocessing
data_augmentation = keras.Sequential([
```

```

        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.08),
        layers.RandomZoom(0.08),
        layers.RandomTranslation(0.06, 0.06)
    ], name="data_augmentation")

# Visualize augmentation example
plt.figure(figsize=(8,4))
for i in range(6):
    aug_img = data_augmentation(tf.expand_dims(x_train[i], 0), training=True)
    plt.subplot(2,3,i+1)
    plt.imshow(aug_img[0].numpy())
    plt.axis('off')
plt.suptitle("Augmented samples")
plt.show()

```

Augmented samples



```

# Cell 5: Build baseline CNN model (simple and effective)
def create_baseline_model(input_shape=(32,32,3), num_classes=10):
    inputs = layers.Input(shape=input_shape)
    x = data_augmentation(inputs) # include augmentation in model so it's applied during training
    x = layers.Conv2D(32, (3,3), padding='same', activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Conv2D(32, (3,3), padding='same', activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.MaxPooling2D((2,2))(x)
    x = layers.Dropout(0.25)(x)

    x = layers.Conv2D(64, (3,3), padding='same', activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Conv2D(64, (3,3), padding='same', activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.MaxPooling2D((2,2))(x)
    x = layers.Dropout(0.3)(x)

    x = layers.Conv2D(128, (3,3), padding='same', activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.MaxPooling2D((2,2))(x)
    x = layers.Dropout(0.4)(x)

    x = layers.Flatten()(x)
    x = layers.Dense(256, activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)

    model = keras.Model(inputs, outputs, name='cifar10_cnn_baseline')
    return model

model = create_baseline_model()
model.summary()

```

Model: "cifar10_cnn_baseline"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 32, 32, 3)	0
data_augmentation (Sequential)	(None, 32, 32, 3)	0
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (BatchNormalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None, 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_2 (Dropout)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 256)	524,544
batch_normalization_5 (BatchNormalization)	(None, 256)	1,024
dropout_3 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 10)	2,570

```
# Cell 6: Compile and callbacks
batch_size = 64
epochs = 50

model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-3),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

checkpoint_cb = callbacks.ModelCheckpoint("best_cifar10_cnn.h5", save_best_only=True, monitor='val_accuracy', mode='max')
earlystop_cb = callbacks.EarlyStopping(monitor='val_loss', patience=8, restore_best_weights=True)
reduce_lr_cb = callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6)
```

```
# Cell 7: Train the model
history = model.fit(
    x_train, y_train_flat,
    epochs=epochs,
    batch_size=batch_size,
    validation_split=0.1,
    callbacks=[checkpoint_cb, earlystop_cb, reduce_lr_cb],
    shuffle=True
```

)

```

Epoch 1/50
704/704 0s 19ms/step - accuracy: 0.3003 - loss: 2.2442WARNING:absl:You are saving your model as an HDF5 f
704/704 23s 20ms/step - accuracy: 0.3004 - loss: 2.2437 - val_accuracy: 0.4912 - val_loss: 1.3881 - learn
Epoch 2/50
703/704 0s 19ms/step - accuracy: 0.4805 - loss: 1.4517WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.4806 - loss: 1.4516 - val_accuracy: 0.5600 - val_loss: 1.2170 - learn
Epoch 3/50
704/704 13s 19ms/step - accuracy: 0.5390 - loss: 1.2852 - val_accuracy: 0.5372 - val_loss: 1.3694 - learn
Epoch 4/50
703/704 0s 20ms/step - accuracy: 0.5832 - loss: 1.1664WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 21ms/step - accuracy: 0.5832 - loss: 1.1663 - val_accuracy: 0.5820 - val_loss: 1.3885 - learn
Epoch 5/50
701/704 0s 19ms/step - accuracy: 0.6118 - loss: 1.0942WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 20ms/step - accuracy: 0.6118 - loss: 1.0941 - val_accuracy: 0.6470 - val_loss: 1.0259 - learn
Epoch 6/50
704/704 14s 19ms/step - accuracy: 0.6320 - loss: 1.0398 - val_accuracy: 0.6404 - val_loss: 1.0445 - learn
Epoch 7/50
704/704 14s 19ms/step - accuracy: 0.6434 - loss: 1.0154 - val_accuracy: 0.5682 - val_loss: 1.3727 - learn
Epoch 8/50
704/704 14s 19ms/step - accuracy: 0.6690 - loss: 0.9554 - val_accuracy: 0.6406 - val_loss: 1.0825 - learn
Epoch 9/50
703/704 0s 19ms/step - accuracy: 0.6841 - loss: 0.9050WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.6841 - loss: 0.9050 - val_accuracy: 0.6596 - val_loss: 1.0172 - learn
Epoch 10/50
701/704 0s 19ms/step - accuracy: 0.6923 - loss: 0.8798WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.6923 - loss: 0.8798 - val_accuracy: 0.6764 - val_loss: 0.9405 - learn
Epoch 11/50
703/704 0s 19ms/step - accuracy: 0.6994 - loss: 0.8600WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.6994 - loss: 0.8600 - val_accuracy: 0.7074 - val_loss: 0.8439 - learn
Epoch 12/50
704/704 14s 20ms/step - accuracy: 0.7044 - loss: 0.8500 - val_accuracy: 0.6828 - val_loss: 0.9446 - learn
Epoch 13/50
703/704 0s 19ms/step - accuracy: 0.7118 - loss: 0.8253WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.7118 - loss: 0.8253 - val_accuracy: 0.7386 - val_loss: 0.7604 - learn
Epoch 14/50
704/704 14s 19ms/step - accuracy: 0.7180 - loss: 0.8123 - val_accuracy: 0.6862 - val_loss: 0.9092 - learn
Epoch 15/50
704/704 13s 19ms/step - accuracy: 0.7215 - loss: 0.8043 - val_accuracy: 0.7312 - val_loss: 0.7876 - learn
Epoch 16/50
702/704 0s 19ms/step - accuracy: 0.7256 - loss: 0.7905WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.7256 - loss: 0.7905 - val_accuracy: 0.7490 - val_loss: 0.7173 - learn
Epoch 17/50
704/704 14s 19ms/step - accuracy: 0.7298 - loss: 0.7832 - val_accuracy: 0.7406 - val_loss: 0.7936 - learn
Epoch 18/50
702/704 0s 19ms/step - accuracy: 0.7349 - loss: 0.7719WARNING:absl:You are saving your model as an HDF5 f
704/704 14s 19ms/step - accuracy: 0.7349 - loss: 0.7720 - val_accuracy: 0.7686 - val_loss: 0.6665 - learn
Epoch 19/50
704/704 14s 19ms/step - accuracy: 0.7302 - loss: 0.7779 - val_accuracy: 0.7364 - val_loss: 0.7648 - learn
Epoch 20/50
704/704 14s 20ms/step - accuracy: 0.7385 - loss: 0.7517 - val_accuracy: 0.7574 - val_loss: 0.7175 - learn
Epoch 21/50
704/704 14s 19ms/step - accuracy: 0.7387 - loss: 0.7499 - val_accuracy: 0.7114 - val_loss: 0.8673 - learn
Epoch 22/50
704/704 14s 19ms/step - accuracy: 0.7500 - loss: 0.7197 - val_accuracy: 0.7568 - val_loss: 0.7222 - learn
Epoch 23/50
704/704 14s 19ms/step - accuracy: 0.7506 - loss: 0.7166 - val_accuracy: 0.7556 - val_loss: 0.7198 - learn
Epoch 24/50
    ... 100 / +           0.7500  1   0.7076  7   0.7200  1   0.7198  7


```

```

# Cell 8: Plot training curves
plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['loss'], label='train_loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.legend()
plt.title('Loss')

plt.subplot(1,2,2)
plt.plot(history.history['accuracy'], label='train_acc')
plt.plot(history.history['val_accuracy'], label='val_acc')
plt.legend()
plt.title('Accuracy')
plt.show()

```




```
# Cell 9: Evaluate on test set
test_loss, test_acc = model.evaluate(x_test, y_test_flat, verbose=2)
print(f"Test loss: {test_loss:.4f}, Test accuracy: {test_acc:.4f}")

# Predictions and confusion matrix
y_pred_probs = model.predict(x_test)
y_pred = np.argmax(y_pred_probs, axis=1)

cm = confusion_matrix(y_test_flat, y_pred)
plt.figure(figsize=(10,8))
sns.heatmap(cm, annot=True, fmt='d', xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

print(classification_report(y_test_flat, y_pred, target_names=class_names))
```

313/313 - 3s - 10ms/step - accuracy: 0.7584 - loss: 0.6991
Test loss: 0.6991, Test accuracy: 0.7584
313/313 2s 7ms/step



Predicted

	precision	recall	f1-score	support
airplane	0.72	0.87	0.78	1000
automobile	0.86	0.93	0.89	1000
bird	0.85	0.51	0.64	1000
cat	0.72	0.47	0.57	1000
deer	0.73	0.73	0.73	1000
dog	0.82	0.55	0.66	1000
frog	0.59	0.94	0.73	1000
horse	0.75	0.88	0.81	1000
ship	0.91	0.84	0.87	1000
truck	0.79	0.87	0.83	1000
accuracy			0.76	10000
macro avg	0.77	0.76	0.75	10000
weighted avg	0.77	0.76	0.75	10000

```
# Cell 10: display some predictions
n = 12
indices = np.random.choice(len(x_test), n, replace=False)
plt.figure(figsize=(14,6))
for i, idx in enumerate(indices):
    plt.subplot(3,4,i+1)
    plt.imshow(x_test[idx])
    true_label = class_names[y_test_flat[idx]]
    pred_label = class_names[y_pred[idx]]
    plt.title(f"T:{true_label}\nP:{pred_label}")
    plt.axis('off')
plt.tight_layout()
```

T:cat

T:airplane

T:frog

T:bird